



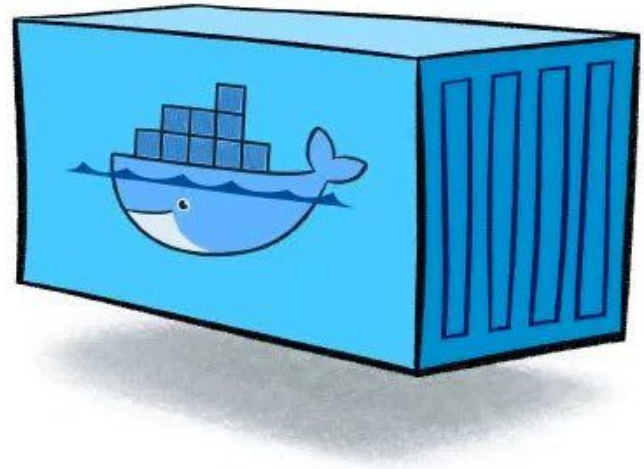
# Module Five

Image Deep Dive



# Image recap

# Images = shipping containers



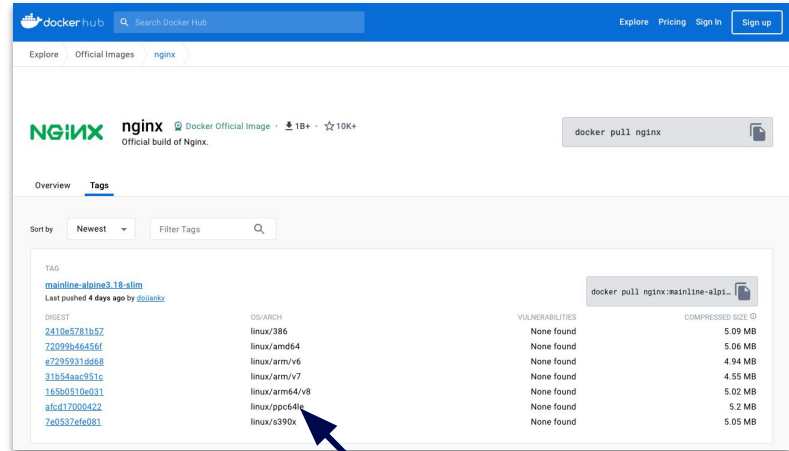
# The container specifications

- Docker created the Open Container Initiative in June 2015
- Currently owned by the Linux Foundation
- Currently defines three specifications
  - **image-spec** - defines image structures and manifests
  - **runtime-spec** - defines how to run OCI images
  - **distribution-spec** - defines the API protocol to push, pull, and discovery content



# Image architectures

- Binaries are compiled to their underlying architectures
- Multi-arch images make it possible to run on multiple architectures
- The container runtime will automatically pull the appropriate architecture



The screenshot shows the Docker Hub interface for the `nginx` image. The page displays the `nginx` image with a pull button. Below the image name, there is a section for tags, showing the `mainline-alpine3.18-slim` tag. A table lists the architectures supported by this image, showing that it is multi-architectured.

DIGEST	OS/ARCH	VULNERABILITIES	COMPRESSED SIZE
2410e5781ba57	linux/386	None found	5.09 MB
72099b46456f	linux/amd64	None found	5.06 MB
e7295931d668	linux/arm/v6	None found	4.94 MB
31b5aac951c	linux/arm/v7	None found	4.55 MB
168b0510e031	linux/arm64/v8	None found	5.02 MB
afcd17000422	linux/ppc64le	None found	5.2 MB
7e0537efe081	linux/s390x	None found	5.05 MB

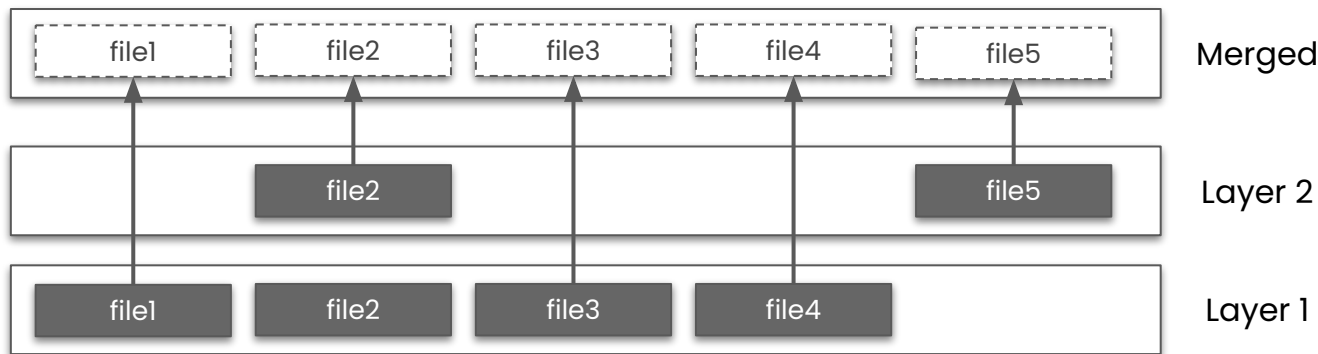




# Understanding unioned filesystems

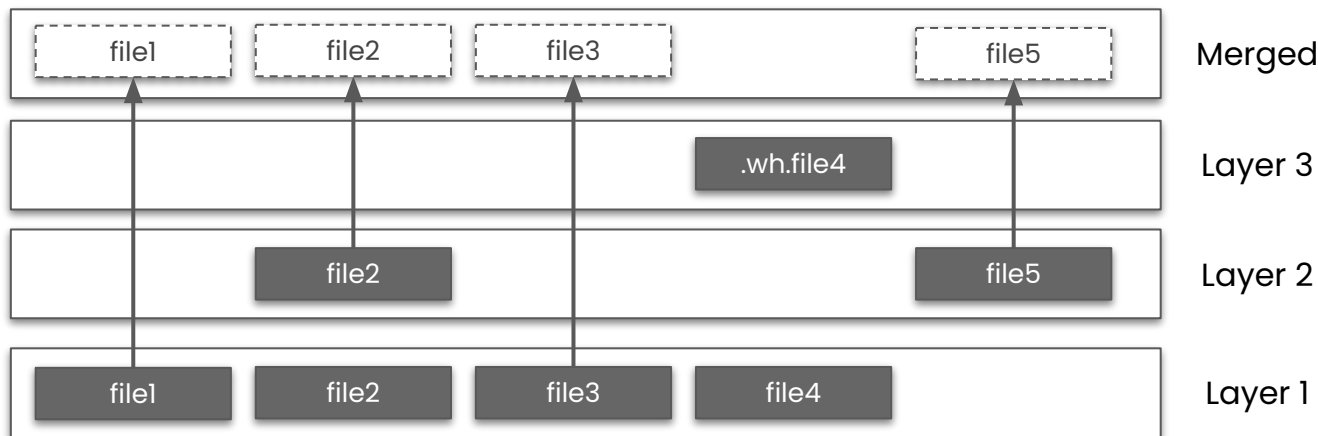
# Image layering

- Unioned filesystems give the ability of combining multiple directories to make a single, unioned filesystem
  - Each layer can add files as needed
  - Files in “higher” layers replace those from “lower” layers



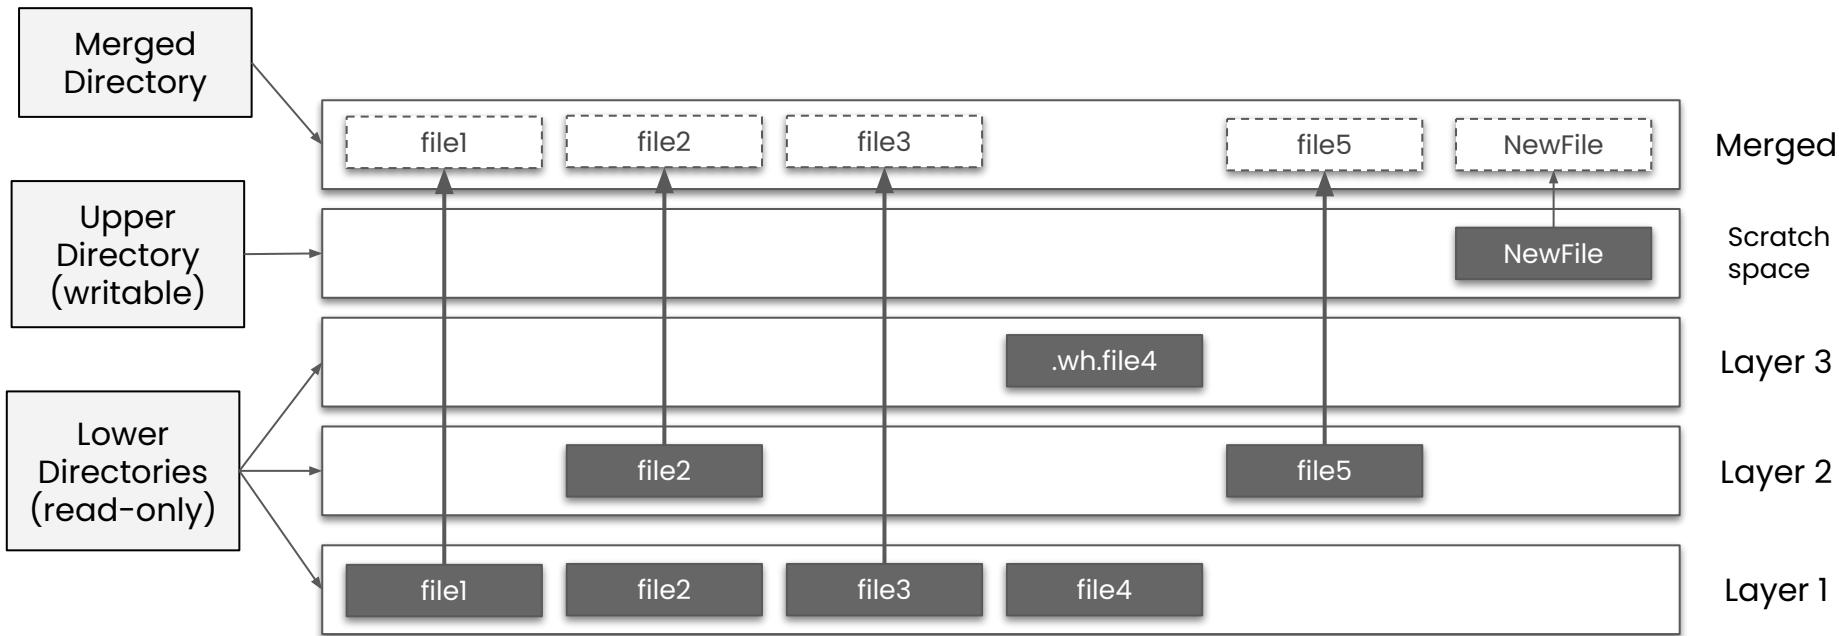
# Deleted files

- If a layer wants to delete a file, it creates a “whiteout” file
  - Whiteout files are only used by the filesystem driver and not visible in the merged filesystem





# Union filesystem terminology



All of this is viewable using `docker inspect <container>`



# Creating an image manually

- Images can be manually created using these steps:
  - **docker run** - start a new container
  - **docker exec** and **docker cp** to make changes
  - **docker commit** to save the container's filesystem as an image



# Or use a Dockerfile!

- Most often included with source code repositories
- Text-based file with instructions on how to build an image
  - **FROM** – the base to start from
  - **WORKDIR** – set the working directory in the image
  - **COPY** – copy files from host into the image
  - **RUN** – run a command
  - **CMD** – set the default command
- Build it with **docker build**

```
FROM ubuntu
WORKDIR /usr/local/app
RUN apt update && apt install -y nodejs
COPY index.js .
CMD ["node", "index.js"]
```



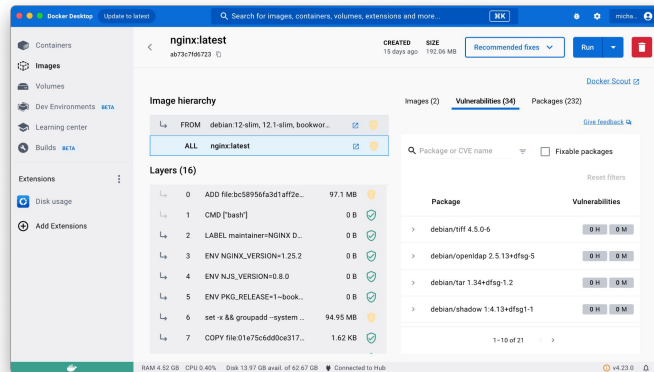
# Visualizing the layers

- **docker image history** – view details about each layer
- Image analysis in Docker Desktop
- Other open-source tools
  - [github.com/wagoodman/dive](https://github.com/wagoodman/dive) – dive into an image and see changes

```
-zsh

> docker image history nginx
IMAGE          CREATED          CREATED BY          SIZE      COMMENT
ab73c7fd6723   2 weeks ago     /bin/sh -c #(nop)  CMD ["nginx" "-g" "daemon_  0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  STOP SIGNAL SIGQUIT  0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  EXPOSE 80 0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  ENTRYPOINT ["/docker-entr_  0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  COPY file:9s3b2b63db9f8fc7_  4.62kB
<missing>      2 weeks ago     /bin/sh -c #(nop)  COPY file:57846632acc8975_  3.82kB
<missing>      2 weeks ago     /bin/sh -c #(nop)  COPY file:3b1b9915b7dd898a_  298B
<missing>      2 weeks ago     /bin/sh -c #(nop)  COPY file:caec368f5a54f70a_  2.12kB
<missing>      2 weeks ago     /bin/sh -c #(nop)  COPY file:01e75c6dd0ce317d_  1.62kB
<missing>      2 weeks ago     /bin/sh -c #(nop)  set -x 66 groupadd --system --  94.9MB
<missing>      2 weeks ago     /bin/sh -c #(nop)  ENV PKG_RELEASE=1~bookworm 0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  ENV NJ5_VERSION=0.8.0 0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  ENV NGINX_VERSION=1.25.2 0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  LABEL maintainer=NGINX Do... 0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  CMD ["bash"] 0B
<missing>      2 weeks ago     /bin/sh -c #(nop)  ADD file:bc58956fa3d1aff2e... 97.1MB

~ |
```

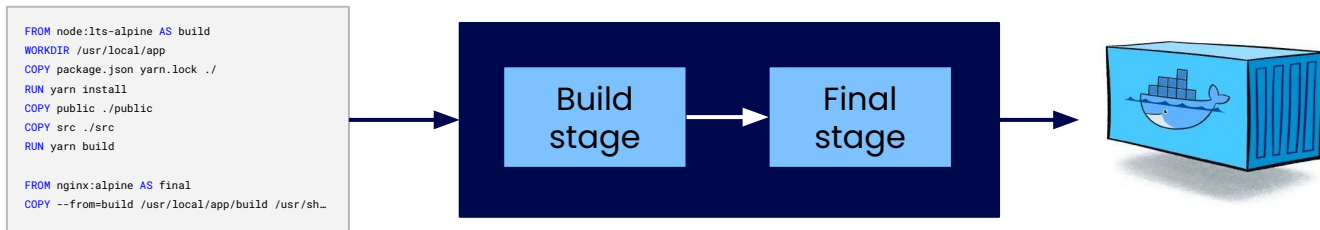




# Overview of builders and terminology

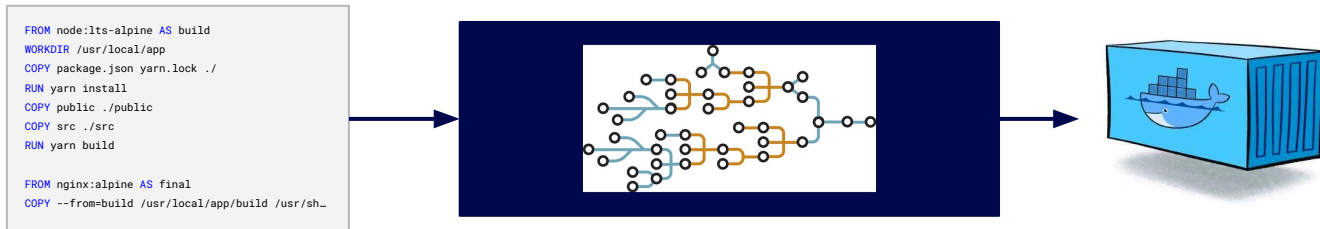
# Legacy Builder

- The original builder used when running **docker build**
- Reads a **Dockerfile** and produces an image
- Very basic support for multi-stage builds
  - Runs all instructions in serial order
  - Builds all stages up to the target stage
- Not able to build multi-architecture images



# BuildKit

- Open-source builder started in Nov 2017
  - Replaced legacy builder in Docker Engine 23.0 (DD 4.19)
- Provides support for caching, distributed workers, new frontends, and much more
- Provides very low-level build definition format
  - You'll most likely never interact with BuildKit directly, but tools that use it



# docker buildx

- A new Docker CLI command that exposes BuildKit features
  - Manages the BuildKit builders
  - Multi-architecture builds
  - Advanced cache management
  - SBOM and provenance creation

```
root@07c6bffa8fee: /  
> docker buildx --help  
  
Usage:  docker buildx [OPTIONS] COMMAND  
  
Extended build capabilities with BuildKit  
  
Options:  
  --builder string  Override the configured builder instance  
  
Management Commands:  
  imagetools       Commands to work on images in registry  
  
Commands:  
  bake             Build from a file  
  build            Start a build  
  create           Create a new builder instance  
  du               Disk usage  
  inspect          Inspect current builder instance  
  ls               List builder instances  
  prune            Remove build cache  
  rm               Remove a builder instance  
  stop             Stop builder instance  
  use              Set the current builder instance  
  version          Show buildx version information  
  
Run 'docker buildx COMMAND --help' for more information on a command.
```







# Using buildx

# Built-in BuildKit

- Docker Desktop includes a BuildKit daemon
- `docker build` has been aliased to `docker buildx build` since DD 4.19
- It has limitations though...
  - Not able to produce multi-architecture images
  - The move to use containerd for image storage should fix this (currently an experimental feature)



# Creating and using builders

- A builder = a BuildKit daemon
- `docker buildx create --name my-builder`
  - Creates a builder named my-builder that will run in a container
- `docker buildx create --platform linux/amd64,linux/arm64`
  - Creates a multi-arch builder that will be given a generated name
- `docker buildx ls`
  - List all of the builders
- `docker buildx use <builder-name>`
  - Make the specified builder the default



# Build drivers

- Drivers provide different ways for how the BuildKit backend runs
- They include:
  - **docker** - uses the BuildKit library bundled into the Docker daemon
  - **docker-container** - creates a dedicated BuildKit container
  - **kubernetes** - creates BuildKit pods in a Kubernetes cluster
  - **remote** - connects directly to a manually managed BuildKit daemon

Feature	docker	docker-container	kubernetes	remote
Automatically load image	✓			
Cache export	Inline only	✓	✓	✓
Tarball output		✓	✓	✓
Multi-arch images		✓	✓	✓
BuildKit configuration		✓	✓	Managed externally



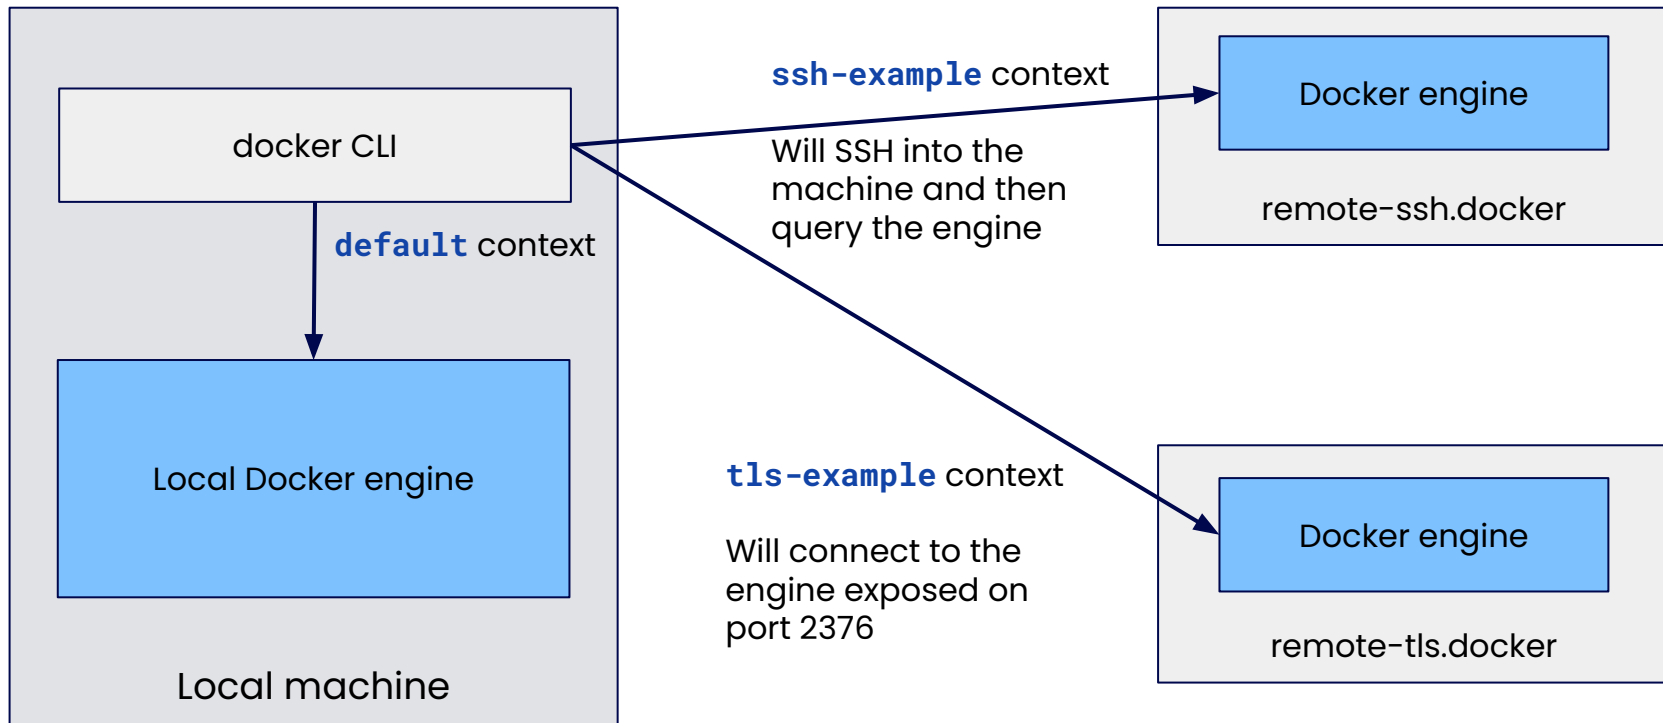
# Docker contexts 101

- Contexts allow your local CLI to work with other Docker engines
- The default context works with the engine on your machine
- Common commands:
  - `docker context create` - create a new context
  - `docker context ls` - list current contexts
  - `docker context use <name>` - make the specified context the default

```
docker context create ssh-example --docker "host=ssh://user@remote-ssh.docker"  
docker context create tls-example \  
--docker "host=tcp://remote-tls.docker:2376,ca=ca.pem,cert=cert.pem,key=key.pem"
```

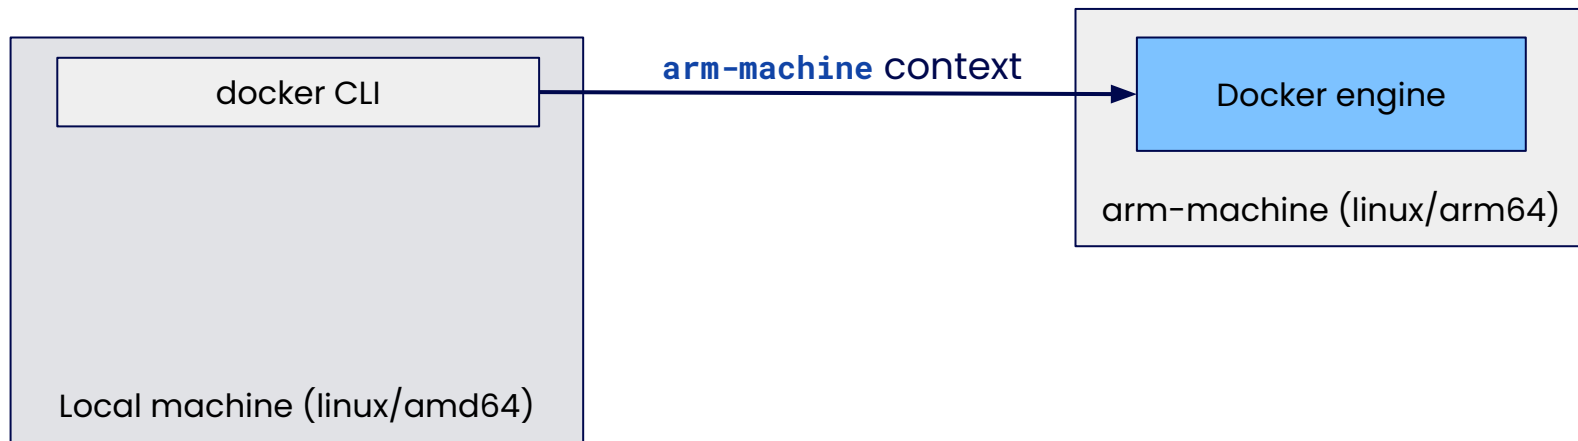


# Docker contexts visually



# Native multi-arch builders

1. First, create a Docker context for the remote nodes
  - `docker context create arm-machine --docker "host=ssh://ubuntu@arm-machine"`
2. Create a builder, specifying the platform and context to use
  - `docker buildx create --platform linux/arm64 --name native-builds arm-machine`



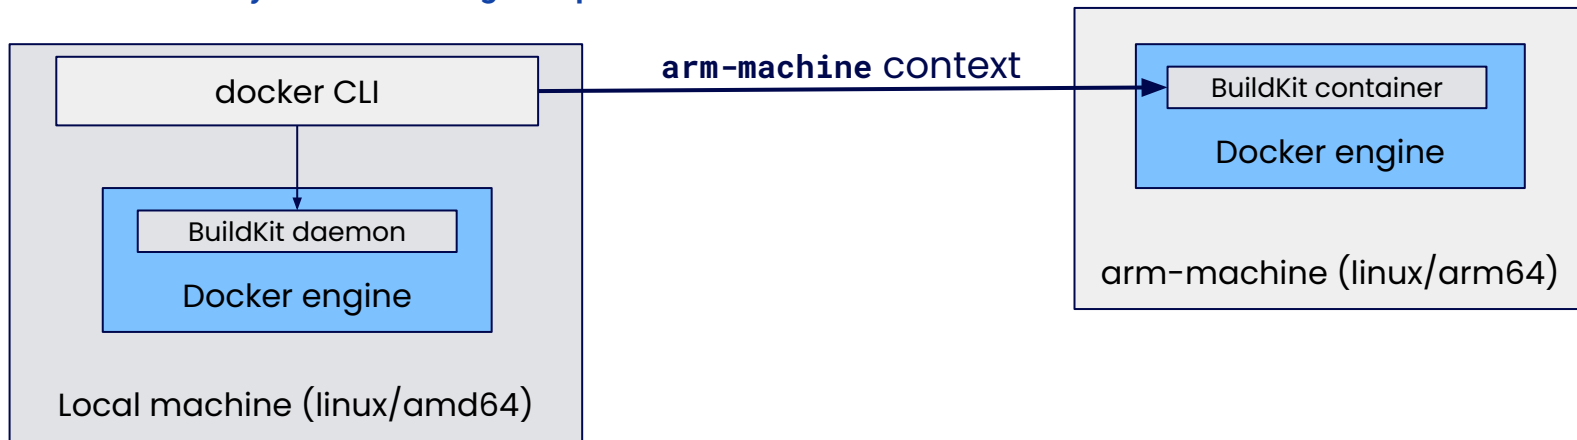
# Native multi-arch builders

## 3. Append another node (our local builder) for amd64 builds

- `docker buildx create --append --name native-builds --platform linux/amd64 default`

## 4. Use the builder to create and push the image

- `docker buildx build --builder native-builds --platform linux/amd64,linux/arm64 \`  
`-t my-custom-image --push .`





# Cache management

- Make the use of various cache backends
  - Great for ephemeral environments or builders (like CI/CD)
- The backends...
  - **inline** - embed the cache into the image itself
  - **registry** - push the cache to a registry
  - **local** - writes the cache to a local directory on the filesystem
  - **gha** - push the build cache to GitHub Actions cache
  - **s3** - uploads the build cache to AWS S3
  - **azblob** - uploads the build cache to Azure Blob Storage
- Cache modes...
  - **min** = only final exported layers are cached (the default)
  - **max** = all intermediate steps are also cached



# Buildx cache examples

- Build and push an image, but store max cache in a registry
  - `docker buildx build --push -t my-repo/my-image \`  
    `--cache-to type=registry,ref=my-repo/my-cache,mode=max \`  
    `--cache-from type=registry,ref=my-repo/my-cache`
- Do the same build, but store the cache in a S3 bucket
  - `docker buildx build --push -t my-repo/my-image \`  
    `--cache-to type=s3,region=us-east-1,bucket=my-cache,prefix=buildx/cache \`  
    `--cache-from type=s3,region=us-east-1,bucket=my-cache,prefix=buildx/cache`



# GitHub Action example

- This action will checkout the code, setup Buildx, login to Hub, build and push the image, and use the GitHub Actions for the build cache

```
jobs:
  docker:
    runs-on: ubuntu-latest
    steps:
      - name: Checkout
        uses: actions/checkout@v4
      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v3
      - name: Login to Docker Hub
        uses: docker/login-action@v3
        with:
          username: ${ secrets.DOCKERHUB_USERNAME }
          password: ${ secrets.DOCKERHUB_TOKEN }
      - name: Build and push
        uses: docker/build-push-action@v5
        with:
          context: .
          push: true
          tags: user/app:latest
          cache-from: type=gha
          cache-to: type=gha,mode=max
```



# Frontends/Syntax

- Buildx supports multiple “frontends” for an image build
  - The frontend converts the file into steps
  - Provides support to test new features
- The frontend is specified using a comment at the top of the file

```
# syntax=docker/dockerfile:1-labs
FROM ubuntu
RUN --security=insecure cat /proc/self/status | grep CapEff
```



# Exporters

- Provide the ability to adjust the output type of a build
- Supported exporters include:
  - **image** - exports to a container image
  - **registry** - exports to a container image and pushes to a registry
  - **local** - exports the build's root filesystem into a local directory
  - **tar** - exports the build's root filesystem into a local tarball
  - **oci** - exports to local filesystem in OCI image layout format
  - **docker** - exports to local filesystem in Docker Image Spec v1.2.0 format
  - **cacheonly** - doesn't export a build, but runs the build and creates cache





# Advanced builds using bake

# Docker Bake

- Higher level build orchestration
- Codify your build commands
- Publish different variants of your images or build several linked projects in parallel
- Uses orchestration file
  - HashiCorp Configuration Language (HCL)
  - JSON
  - YAML (Compose file)



# Complex Docker buildx command without Bake

```
docker buildx build \  
  --push \  
  --cache-from "type=registry,ref=foo/myapp" \  
  --cache-to "type=inline" \  
  --platform "linux/amd64,linux/arm/v6,linux/arm/v7,linux/arm64" \  
  --label "org.opencontainers.image.title=myapp" \  
  --label "org.opencontainers.image.source=https://github.com/foo/myapp" \  
  --label "org.opencontainers.image.version=1.0.0" \  
  --label "org.opencontainers.image.licenses=Apache-2.0" \  
  --tag "foo/myapp:v1.0.0" \  
  --tag "foo/myapp:latest" \  
  --file "./main.Dockerfile" \  
  .
```





# Complex Docker buildx command with Bake

```
{
  target: {
    "docker-metadata-action": {
      output = ["type=registry"]
      cache-from = ["type=registry,ref=foo/myapp"]
      cache-to = ["type=inline"]
      dockerfile = "./main.Dockerfile"
      platforms = ["linux/amd64", "linux/arm/v6", "linux/arm/v7", "linux/arm64"]
      labels: {
        "org.opencontainers.image.title=myapp",
        "org.opencontainers.image.source=https://github.com/foo/myapp",
        "org.opencontainers.image.version=1.0.0",
        "org.opencontainers.image.licenses=Apache-2.0"
      },
      tags: [
        "foo/myapp:v1.0.0",
        "foo/myapp:latest"
      ],
    },
  },
}
```



# Docker Bake – parallel execution

```
group "default" {  
    targets = ["app", "db", "cron"]  
}  
  
target "app" {  
    dockerfile = "Dockerfile.app"  
    platforms = ["linux/amd64", "linux/arm64"]  
    tags = ["repo/app:test"]  
}  
  
target "db" {  
    dockerfile = "Dockerfile.db"  
    platforms = ["linux/amd64", "linux/arm64"]  
    tags = ["repo/db:test"]  
}  
  
target "cron" {  
    dockerfile = "Dockerfile.cron"  
    platforms = ["linux/amd64", "linux/arm64"]  
    tags = ["repo/cron:test"]  
}
```



# Docker Bake – Inherit attributes

```
target "app-dev" {
  args = {
    GO_VERSION = "1.20"
    BUILDX_EXPERIMENTAL = 1
  }
  tags = ["docker.io/username/myapp"]
  dockerfile = "app.Dockerfile"
  labels = {
    "org.opencontainers.image.source" = "https://github.com/username/myapp"
  }
}

target "_release" {
  args = {
    BUILDKIT_CONTEXT_KEEP_GIT_DIR = 1
    BUILDX_EXPERIMENTAL = 0
  }
}

target "app-release" {
  inherits = ["app-dev", "_release"]
  platforms = ["linux/amd64", "linux/arm64"]
}
```



# More Docker Bake Goodness

- Use targets as base images with automatic image dependency resolution
- Variable Block or command line variables
- Call directly from CI/CD tools

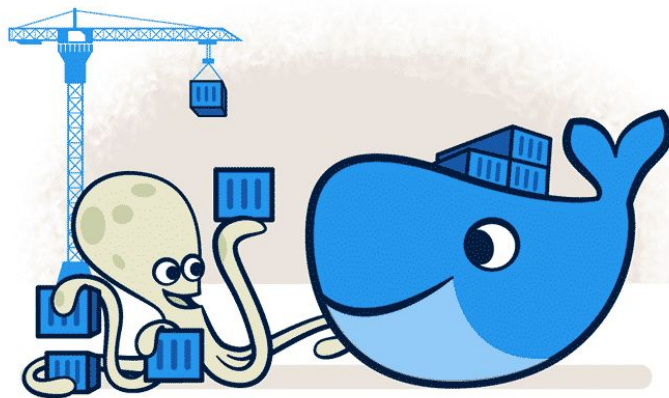
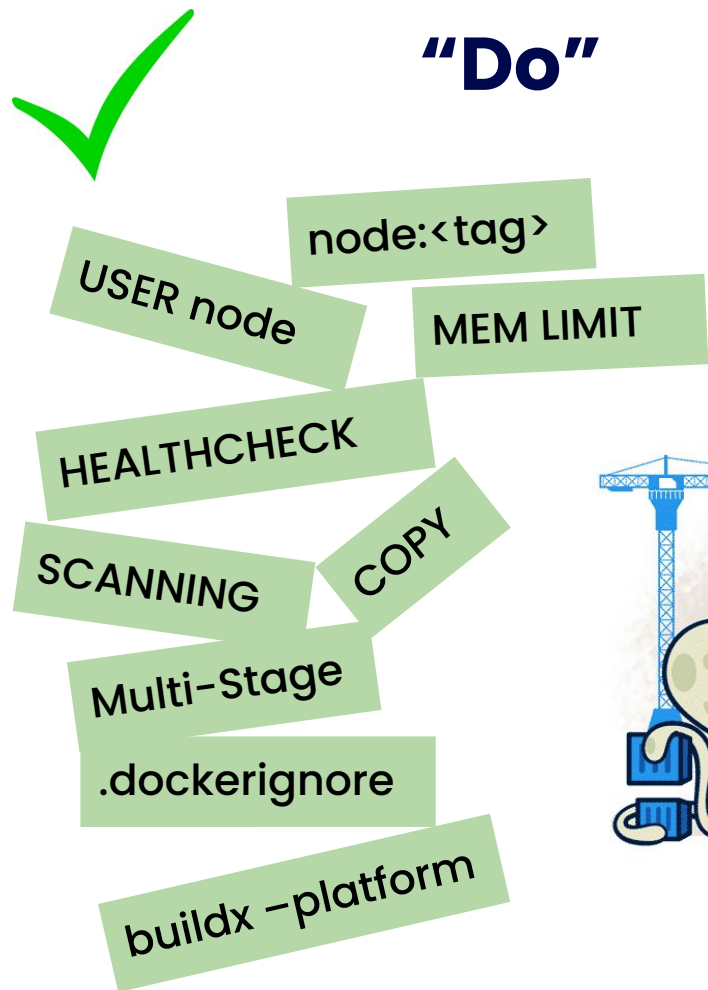




# Best Practices for Images and Builds

**“Do”**

**“Don’t”**





**“Do”**



**“Don’t”**



# Don't mid-air refuel





# Use a vault for secrets



# General Tips

- Each container should do one thing and do it well
  - Avoid creating “general purpose” images
- Don’t ship dev tooling into production
  - Multi-stage image builds help tremendously here!
- Use or create trusted base images
  - Can leverage our Docker Official Images or build your own
- Pin your images (don’t use the default **:latest** tag)
  - **FROM node -> FROM node:20.5-alpine3.17**
  - Can pin to specific SHA sums too! **FROM node@sha256:0b889cbf7...**

Check out the new [“Build with Docker”](#) guide for more build tips!



# Separate app setup from app code

- Install dependencies separate from the app code
- Each of the dependency layers can then be reused

```
FROM node
COPY . .
RUN npm install
EXPOSE 3000
CMD ["node", "src/index.js"]
```



```
FROM node
COPY package.json package-lock.json .
RUN npm install
COPY src ./src
EXPOSE 3000
CMD ["node", "src/index.js"]
```



## Advanced tip - using `--link`

- Indicates a layer is completely independent of other layers
  - Cache busting of previous layers can still allow reuse
- Only works for **COPY** and **ADD** commands

```
FROM node
COPY package.json package-lock.json .
RUN npm install
COPY src ./src
EXPOSE 3000
CMD ["node", "src/index.js"]
```

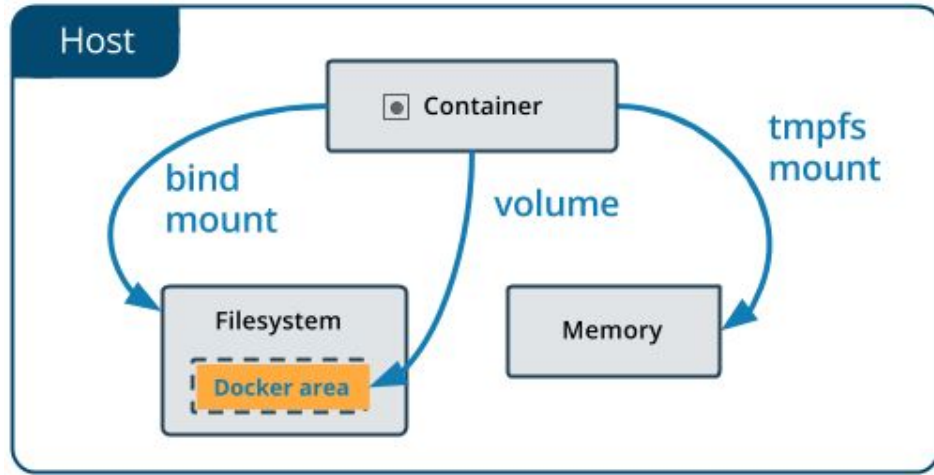


```
FROM node
COPY package.json package-lock.json .
RUN npm install
COPY --link src ./src
EXPOSE 3000
CMD ["node", "src/index.js"]
```



# Volumes

Volumes help store data outside running container. This helps your image stay small and avoids bringing unnecessary or risky data on its journey to production.



# Identifying problems


- Make a small change to your app and watch how many image layers are rebuilt
- In this example, a single file was changed. How many layers are reused?
  - Every code change reshapes dependencies

```
root@ca55e6f6a1ec: /  
> docker build -t node-demo .  
[+] Building 1.6s (9/9) FINISHED                                docker:desktop-linux  
=> [internal] load build definition from Dockerfile              0.0s  
=> => transferring dockerfile: 128B                             0.0s  
=> [internal] load .dockerignore                                0.0s  
=> => transferring context: 2B                                    0.0s  
=> [internal] load metadata for docker.io/library/node:alpine   0.5s  
=> [1/4] FROM docker.io/library/node:alpine@sha256:77516e190b361 0.0s  
=> [internal] load build context                                0.0s  
=> => transferring context: 28.22kB                              0.0s  
=> CACHED [2/4] WORKDIR /usr/local/app                          0.0s  
=> [3/4] COPY . .                                              0.1s  
=> [4/4] RUN npm install                                       0.9s  
=> exporting to image                                           0.0s  
=> => exporting layers                                           0.0s  
=> => writing image sha256:5762d7ddb66912676329cf0ed5af416f5bd0c 0.0s  
=> => naming to docker.io/library/node-demo                    0.0s
```



# Layers

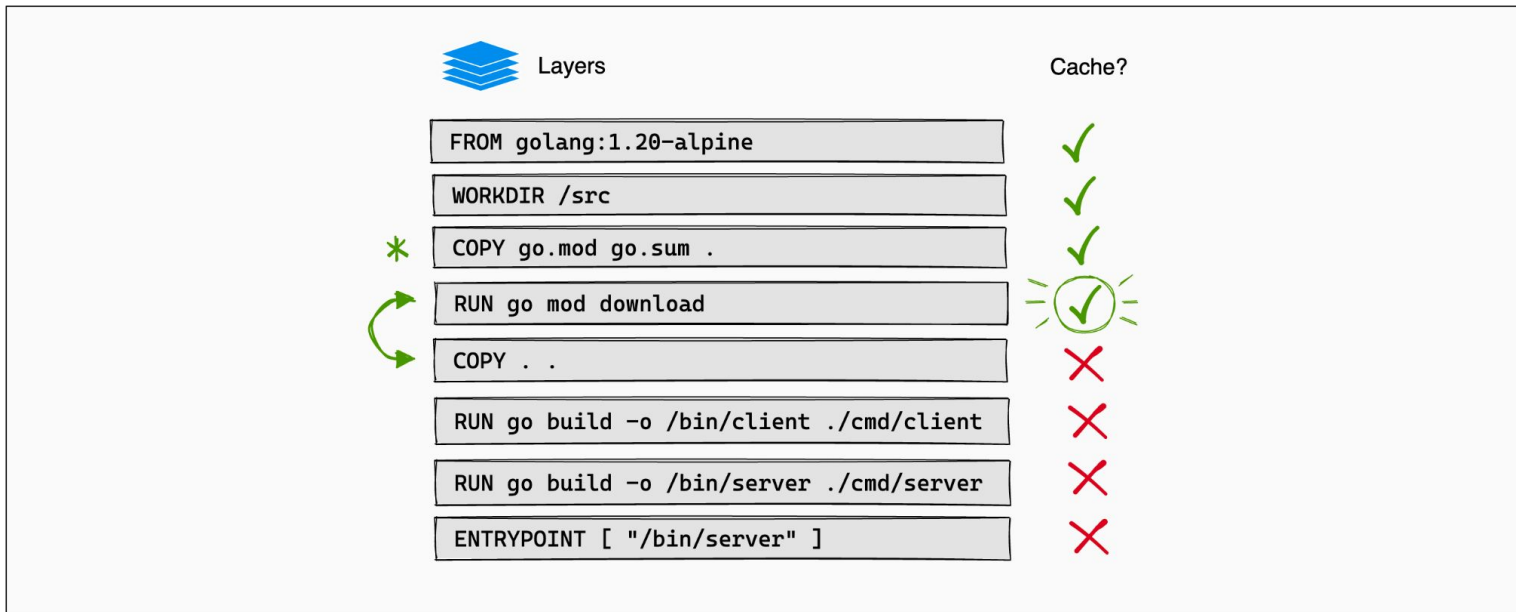
Using layers properly means you're leveraging the build cache.

 Layers	Cache?
FROM golang:1.20-alpine	✓
WORKDIR /src	✓
COPY . .	✗
RUN go mod download	✗
RUN go build -o /bin/client ./cmd/client	✗
RUN go build -o /bin/server ./cmd/server	✗
ENTRYPOINT [ "/bin/server" ]	✗



# Layers

Re-order your commands to better leverage cache.



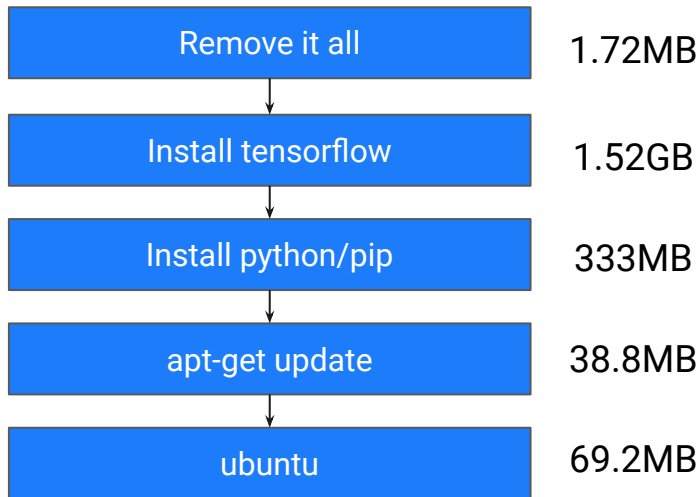


# Looking at a Dockerfile

```
FROM ubuntu
RUN apt update
RUN apt install -y python3 python3-pip
RUN pip install tensorflow
RUN apt autoremove --purge -y python3-pip
```

When we ship this image, we are sending the layer that contains the apt repo cache, the full pip install, and the marker files to indicate we deleted those things.

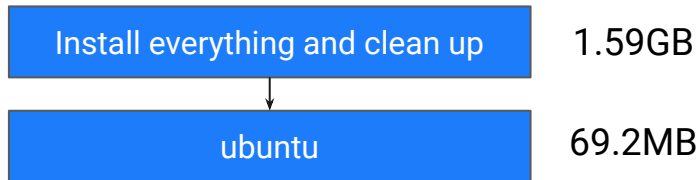
**Tip:** deleting files does NOT affect previous layers, but only what a running container will see



# Cleaning up as we go

- Chain commands into a single **RUN** to combine file changes
- Clean up repo caches after updating/installing

```
FROM ubuntu
RUN apt update && \
    apt install -y python3 python3-pip && \
    pip install tensorflow && \
    apt autoremove --purge -y python3-pip
```



**Net image reduction of 1.96GB to 1.66GB (15% savings)**



# The build cache

- When building an image, Docker tries to reuse layers
- Cache rules...
  - **ADD/COPY** instructions - a checksum of the contents/metadata is created. If the checksum has changed, the cache is invalidated
  - If a **RUN** command changes, the cache is invalidated
  - Any new/removed instructions will invalidate caches
- If the cache is invalidated, that layer is rebuilt, as well as all child layers



# Multi-stage builds

- Multi-stage allows you to create a pipeline within a Dockerfile
- Provides the ability to separate build-time and run-time images
- Use **COPY --from=<stage>** to copy from previous stages
- The last stage is the default target (can be overridden)

```
FROM <image> AS stage1
COPY ...
RUN ...
...

FROM <image2> AS stage2
COPY --from=stage1 /path/in/stage1 /path/in/stage2
```



# A React example

- Use Node to build the React code
- Copy the HTML/CSS/JS into a static web server

```
FROM node:lts AS build
WORKDIR /usr/local/app
COPY package.json package-lock.json ./
RUN npm install
COPY public ./public
COPY src ./src
RUN npm run build
```

```
FROM nginx:alpine
COPY nginx.conf /etc/nginx/nginx.conf
COPY --from=build /usr/local/app/build /usr/share/nginx/html
```



# Using multi-stage in dev

- Create additional stages for dev and target them in Compose
- Stages can start/extend from other stages
- No more need for separate Dockerfiles for dev vs prod!

```
FROM node:lts AS base
WORKDIR /usr/local/app
```

```
FROM base AS dev
CMD ["npm", "run"]
```

```
FROM base AS build
COPY package.json package-lock.json ./
RUN npm install
COPY public ./public
COPY src ./src
RUN npm run build
```

```
FROM nginx:alpine
COPY nginx.conf /etc/nginx/nginx.conf
COPY --from=build /usr/local/app/build /usr/share/nginx/html
```

```
services:
  client:
    build:
      context: .
      target: dev
    volumes:
      - ./usr/local/app
    ports:
      - 3000:3000
```





# Questions and Answers